

Using MySQL for Distributed Database Architectures

Peter Zaitsev

CEO, Percona

PingCAP Infra Meetup,

Shanghai, China, May 26, 2018



About Percona

Solutions for your success with MySQL, MariaDB, MongoDB and PostgreSQL

Support, Managed Services, Consulting, Training, Software

Our Software is 100% Free and Open Source

Support Broad Ecosystem – MySQL, MariaDB, Amazon RDS, Google CloudSQL

In Business for 11 years

More than 3000 customers, including top Internet companies and enterprises

Presentation

Cover Basics

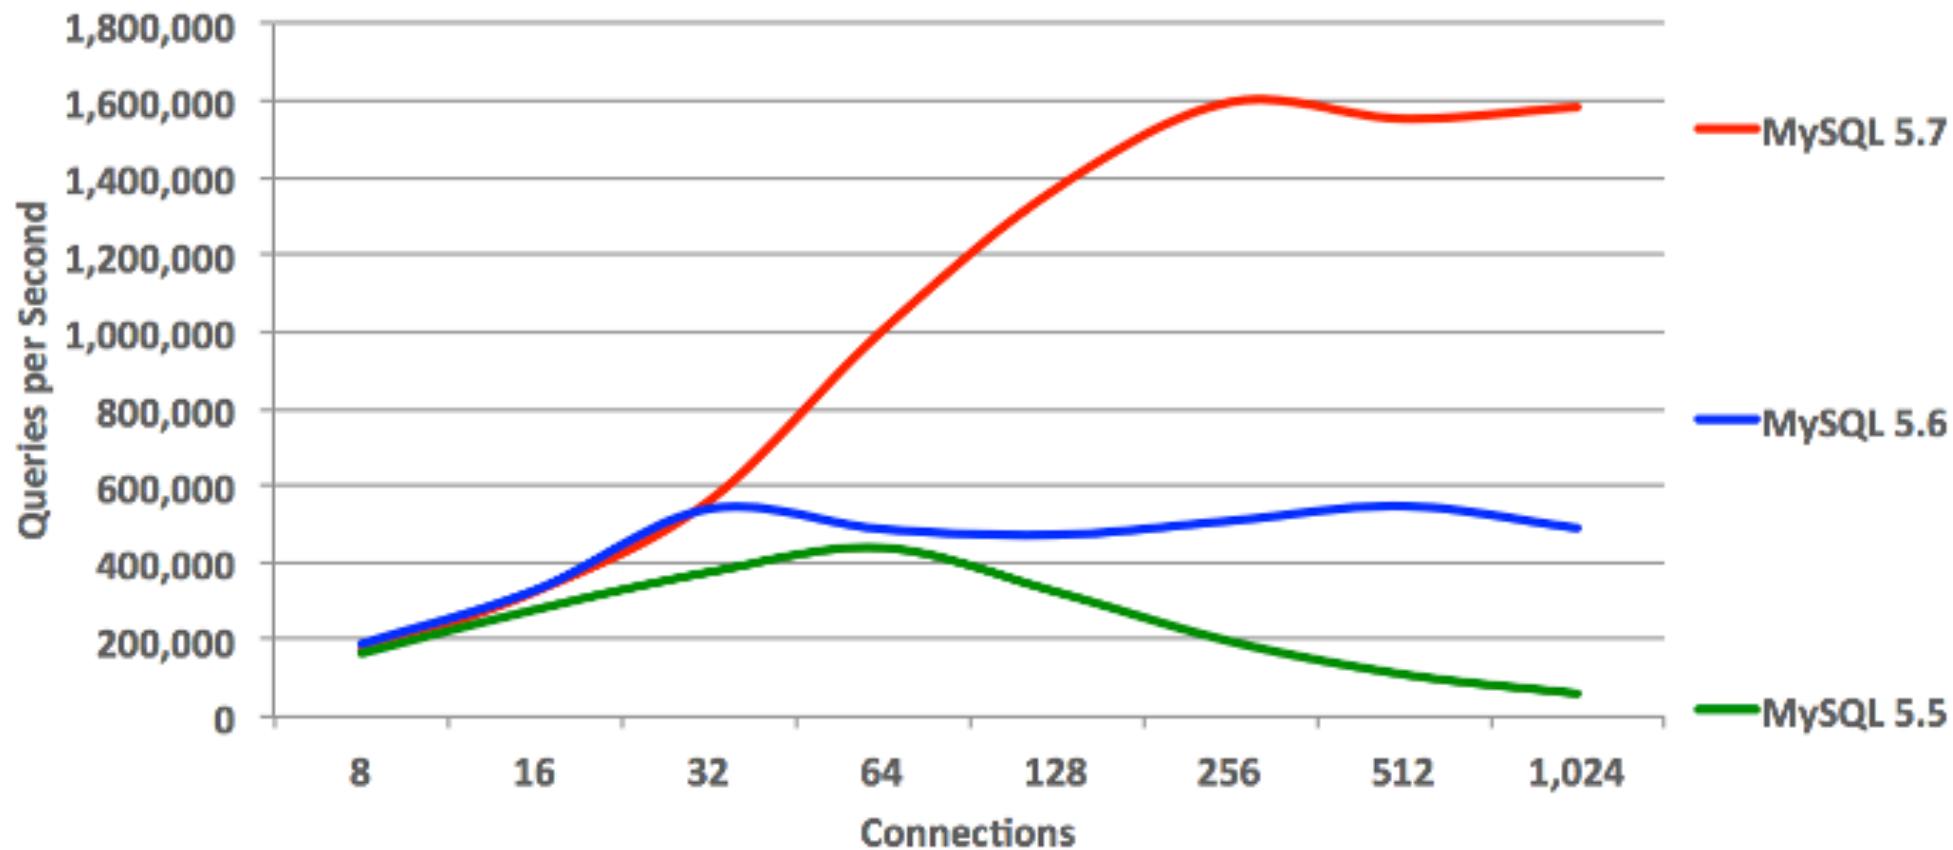
Why Going Distributed

How to do it

Distributed ?

MySQL Deployment on More
than one System

Modern MySQL Scalability



Single MySQL Instance Can Do

Hundreds of Thousands of Queries/Sec

Tends of Thousands of Updates/Sec

Traverse Tens of Millions of Rows/Sec

Comfortably Handle Several TB Database size

Lets Do Some Math

100.000 QPS

10 Queries per User Interaction

10.000 User Interactions/sec

864.000.000 User Interactions/Day

30 User Interactions/User Avg

28.000.000 Daily Active Users Possible

15M of Daily Active Users counting time of day skew

Distributed Systems Tend To be

More Complicated to Develop Against

More Complicated to Operate

Have Additional Performance Bottlenecks

Have Complicated Failure Modes

With All of this ?



Why “Go Distributed” ?

Reasons to “Go Distributed”

High
Availability

Scalability

Data
Distribution

High Availability with MySQL

Cold Standby (ie DRBD)

Failover (Classical Replication)

Active-Active Clustering (PXC, MySQL Group Replication)

Q1:What Failure Modes Do you Consider ?

Server Crash/Server Hardware Failure

Software Bugs and Storage Corruption

Network Failure

“Datacenter” Failure

Developer Mistakes/Intruder Actions

Q2:What Data Am I allowed to lose ?

Limited Data Loss Allowed

No Data Loss

No Transactions In Flight Loss

Q3:How Quickly do you need to Recover ?

“Immediate?”

“Seconds?”

“Minutes?”

“Hours?”

Speed of Light Realities

Data
Propagation
Can Be

- Synchronous – Slow
- Asynchronous –
Data Loss Can Occur

Scalability

Scaling
Reads

Scaling
Writes

Scaling
Data Size

Data Distribution

Some Data
Must be in
Specific
Geographic
Location

- User Latency Reasons
- Legal and Compliance Reasons

Distributed Architectures with MySQL

Main Concepts

Replication

Sharding

Failover Management

Traffic Management

Replication

Having Multiple Copies of the
data, updated with changes

Availability

Service Stays up when
component fails

Availability via Redundancy

Have more than one system

Works well for stateless systems

Is not enough for databases

Availability via Replication

Redundant
Computing
Resource

Paired with
Replicated
Data

Where Replication Happens

Storage Level

Database Level

Application Level

Storage Level Replication

Replication in SAN/NAS, DRBD

Typically provides cold standby

Simple choice which works with many systems

Amazon Aurora – Smart Storage

Database level

Most Flexible

Most Common

Hot/Warm Spare

Some can do Active-Active

Application Level

Hard to get right

Rarely used, especially “used right”

Partial Replication/Synchronization

Smart conflict resolution

Cross Vendor Redundancy

Sharding

Split data set by certain criteria
and store such “shards” on
separate “clusters”

Typical Sharding

By User

By Customer Account/Company

Sharding and Replication

Sharding Almost Always
used with some kind of
Replication

Failover Management

Many Replication Modes Require “Failover”

Need some Tool To Decide on Failover

Need some Tool to Perform Failover

Traffic Management

Scaling with Large Number of Connections

Routing Traffic to Right “Shard”

Read-Write Splitting

Load Management

Avoiding “Dead” Nodes

Options in MySQL

Replication Options in MySQL

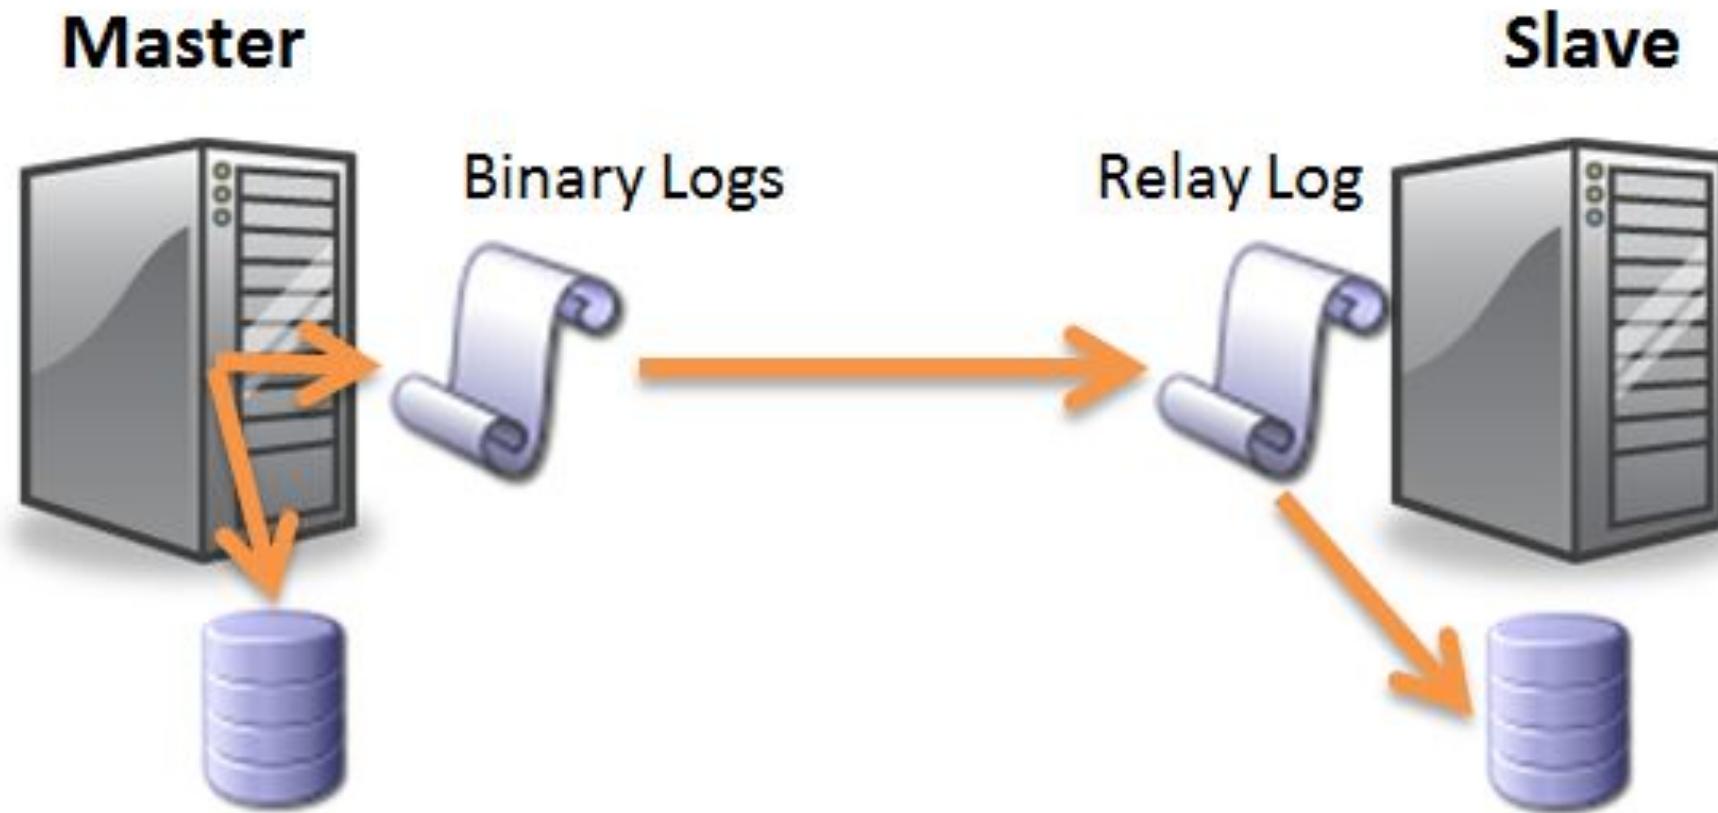
Classical MySQL Replication

MySQL Group Replication

Percona XtraDB Cluster and Galera

MySQL (NDB) Cluster

Classical MySQL Replication



Classical MySQL Replication Properties

Asynchronous or Semi-Synchronous

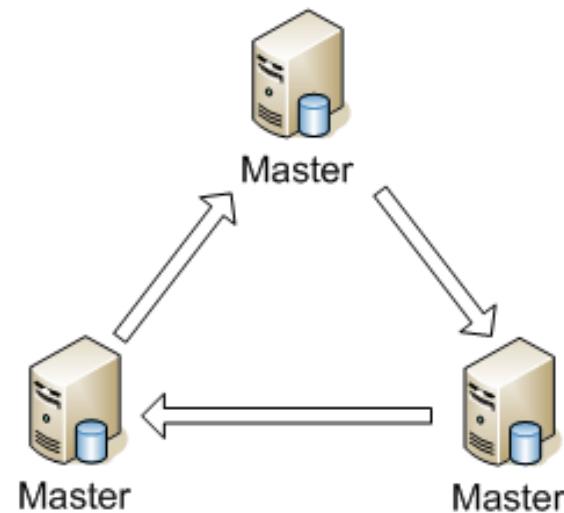
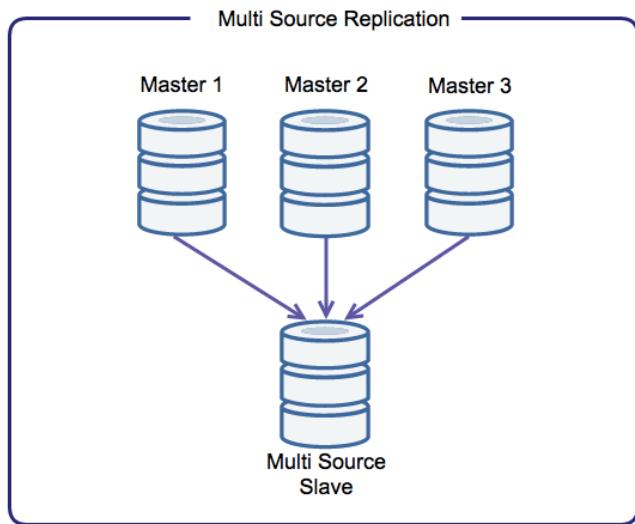
Parallel (since MySQL 5.7)

Many Masters to Many Slaves (MySQL 5.7)

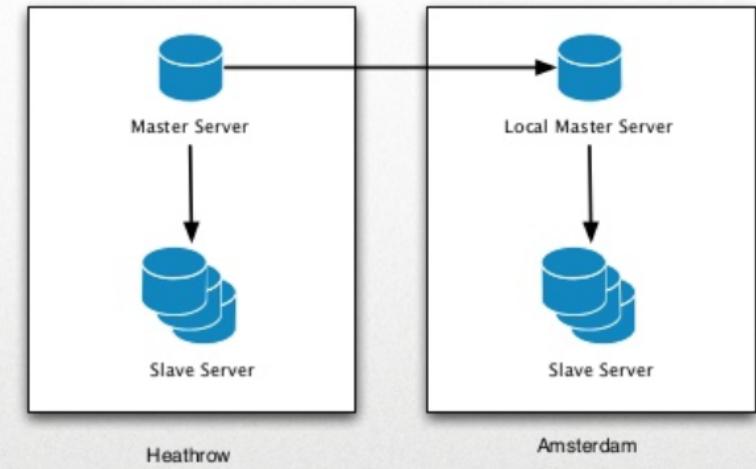
No Conflict Resolution or “Protection”

No Built-in Failover

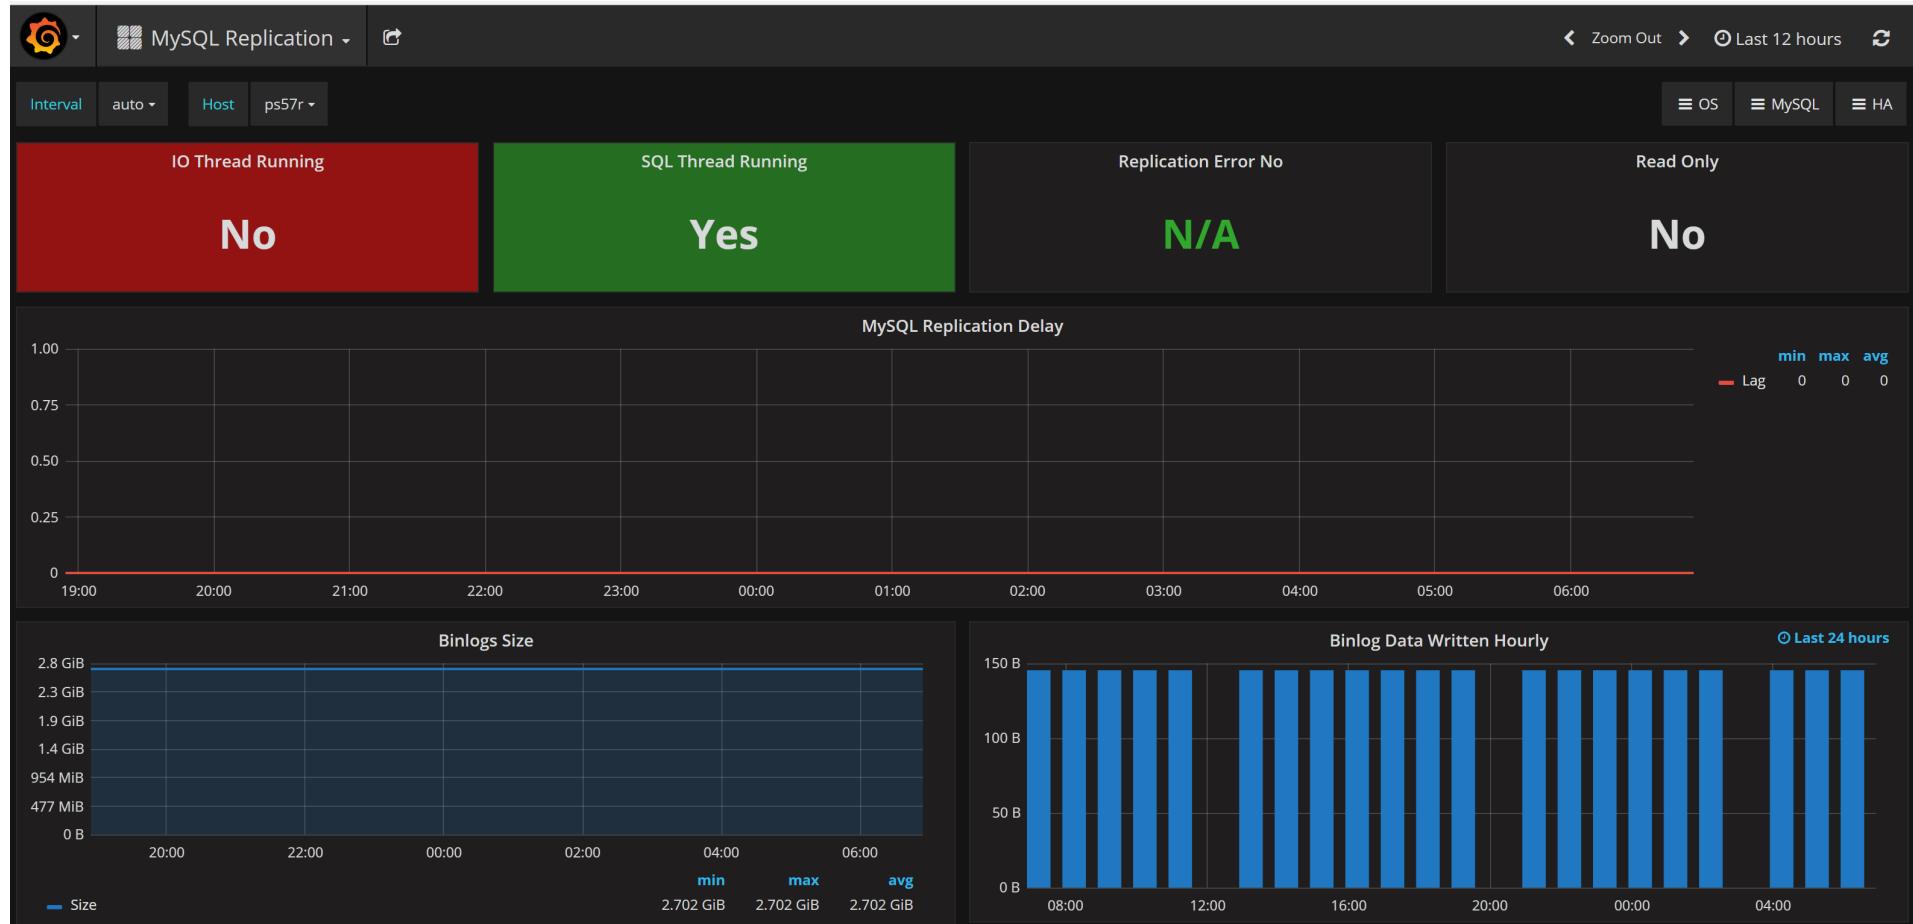
Advanced MySQL Replication Topologies



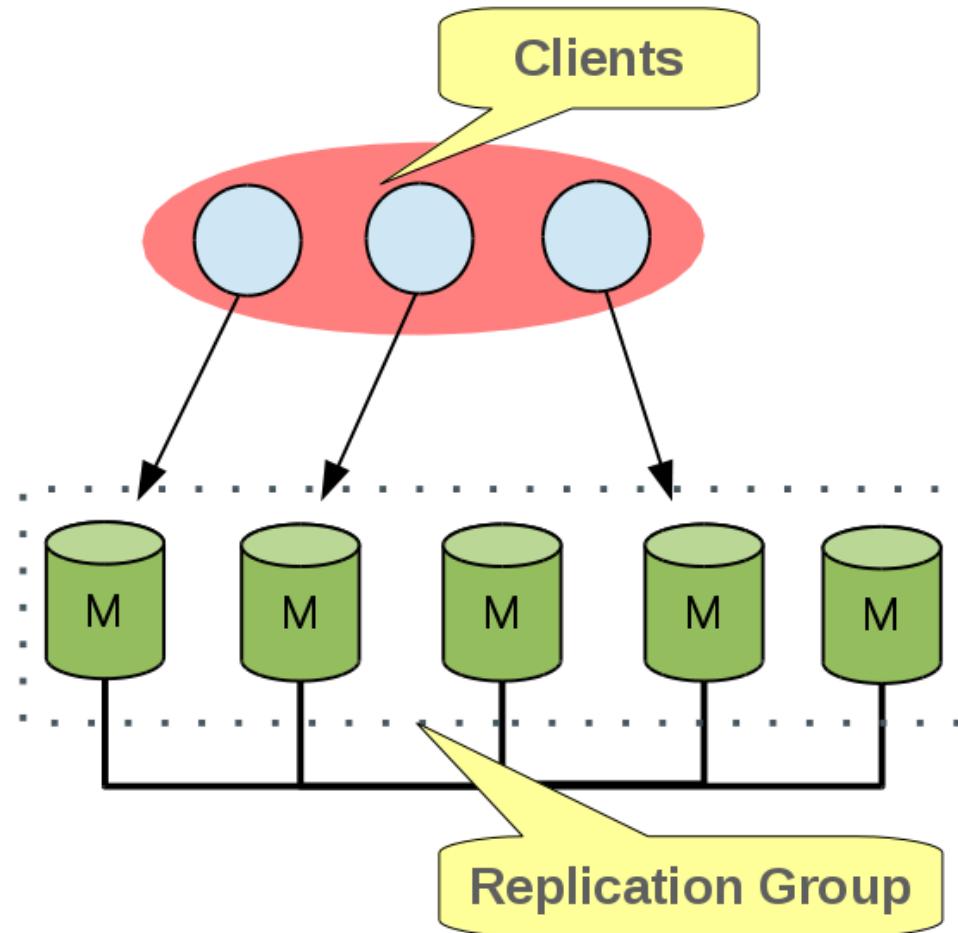
Multi-Level Slaves



PMM Dashboard for Replication



MySQL Group Replication



MySQL Group Replication (New in 5.7)

“Group of Peers”

Write-Anywhere or Dedicated Writer

Asynchronous Replication with Flow Control

Conflicts Prevented through Certification

Built in Failover

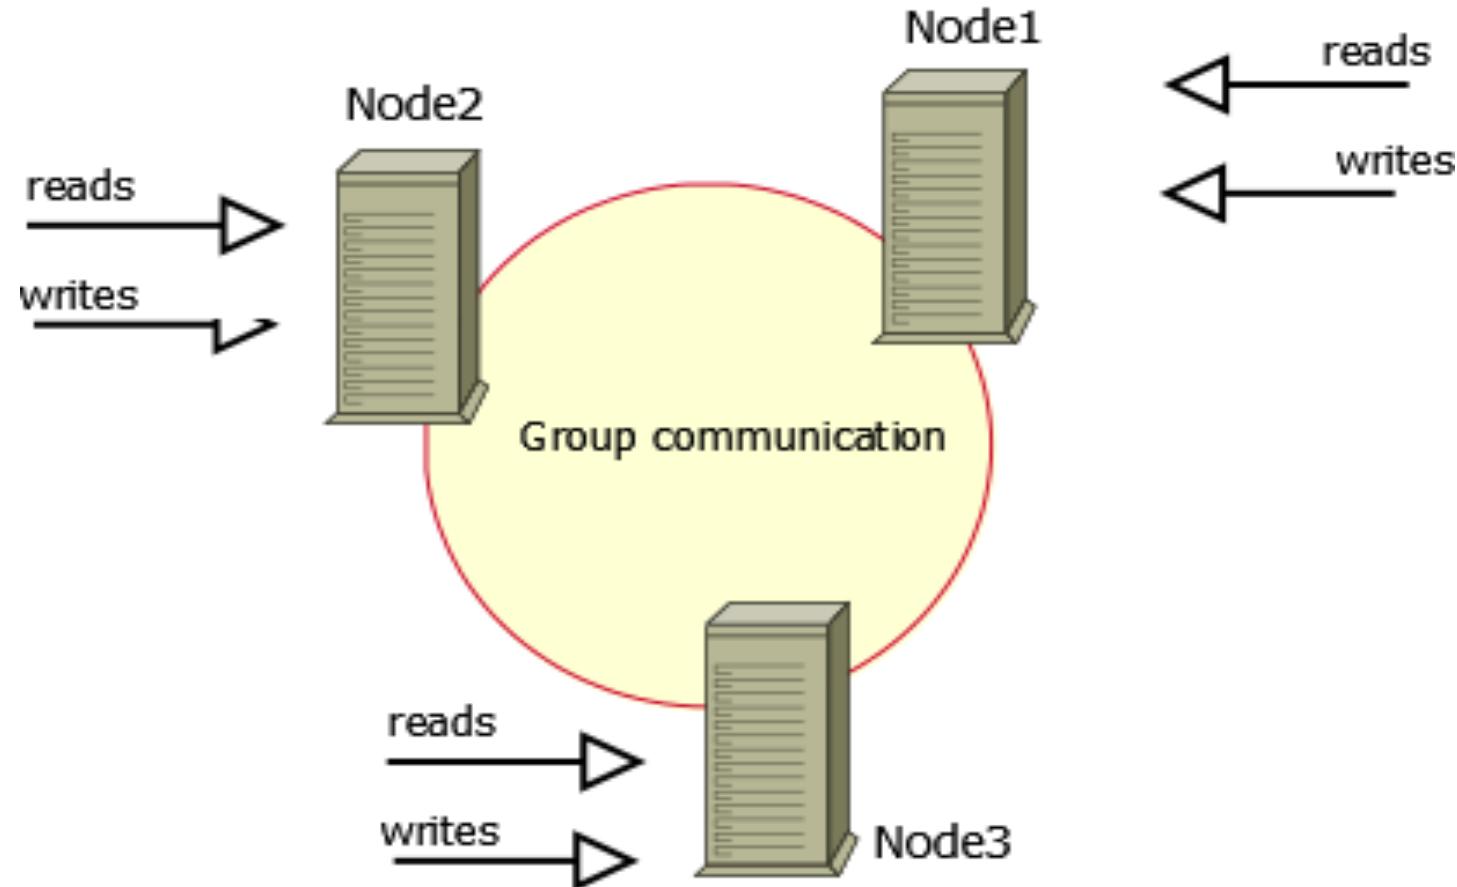
No Automated Provisioning

Percona XtraDB Cluster/Galera



PERCONA
XtraDB Cluster

Percona XtraDB Cluster Topology



PXC/Galera Properties

Write to Any Node

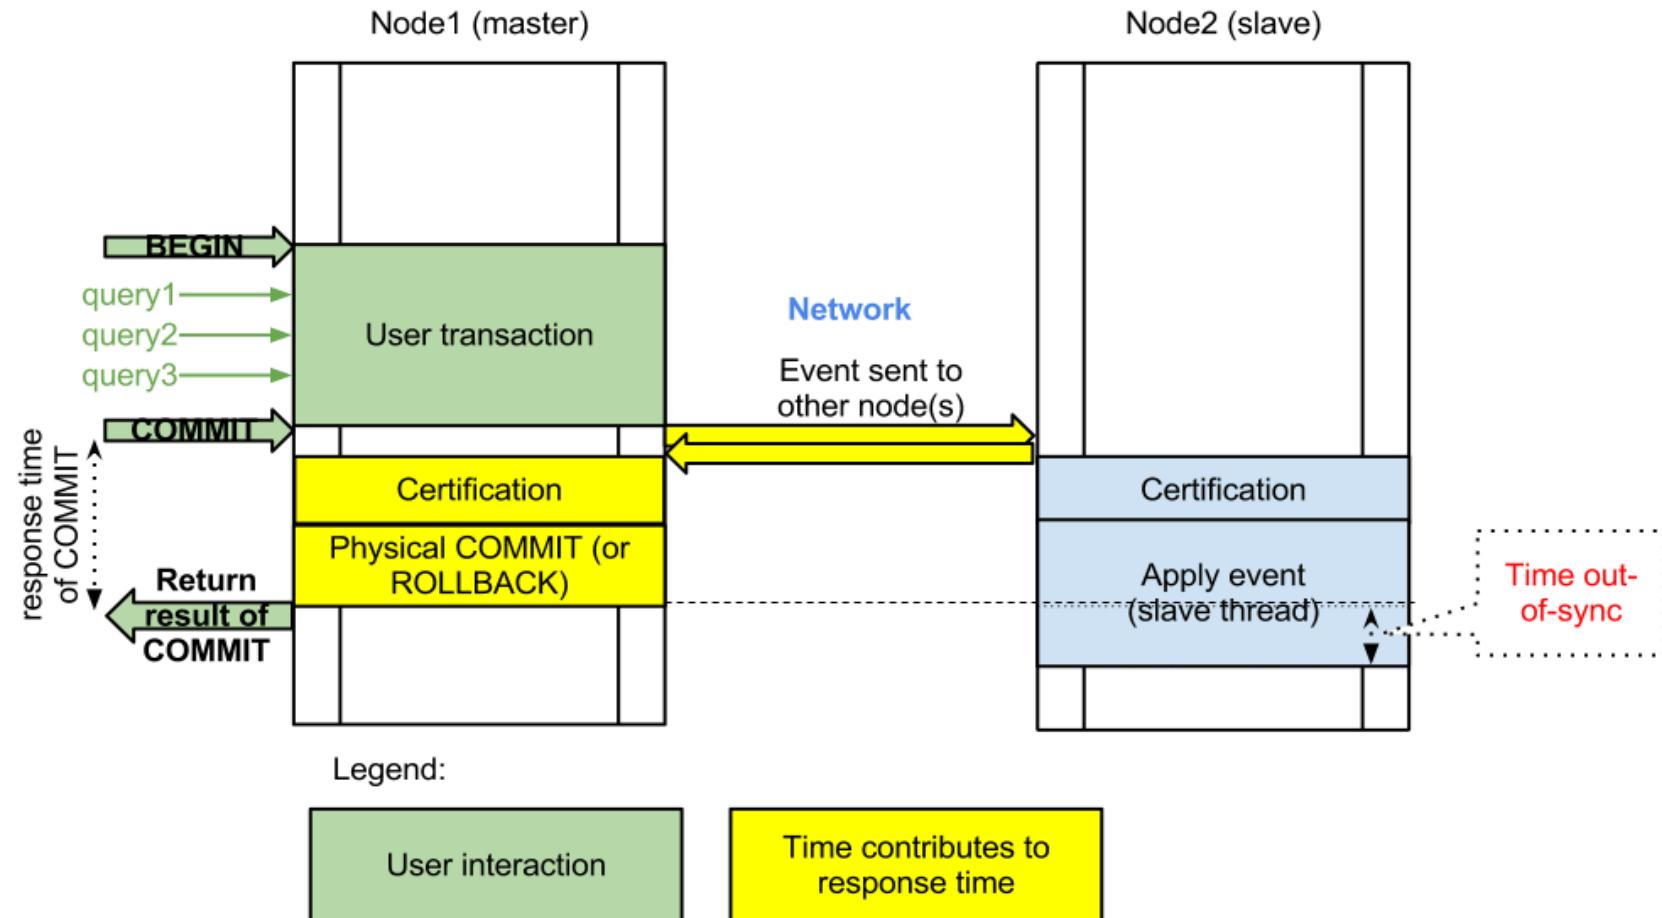
Certification Based Replication

Virtually Synchronous; Can ensure no stale reads

Built in Fail-Over

Built-in Node Provisioning

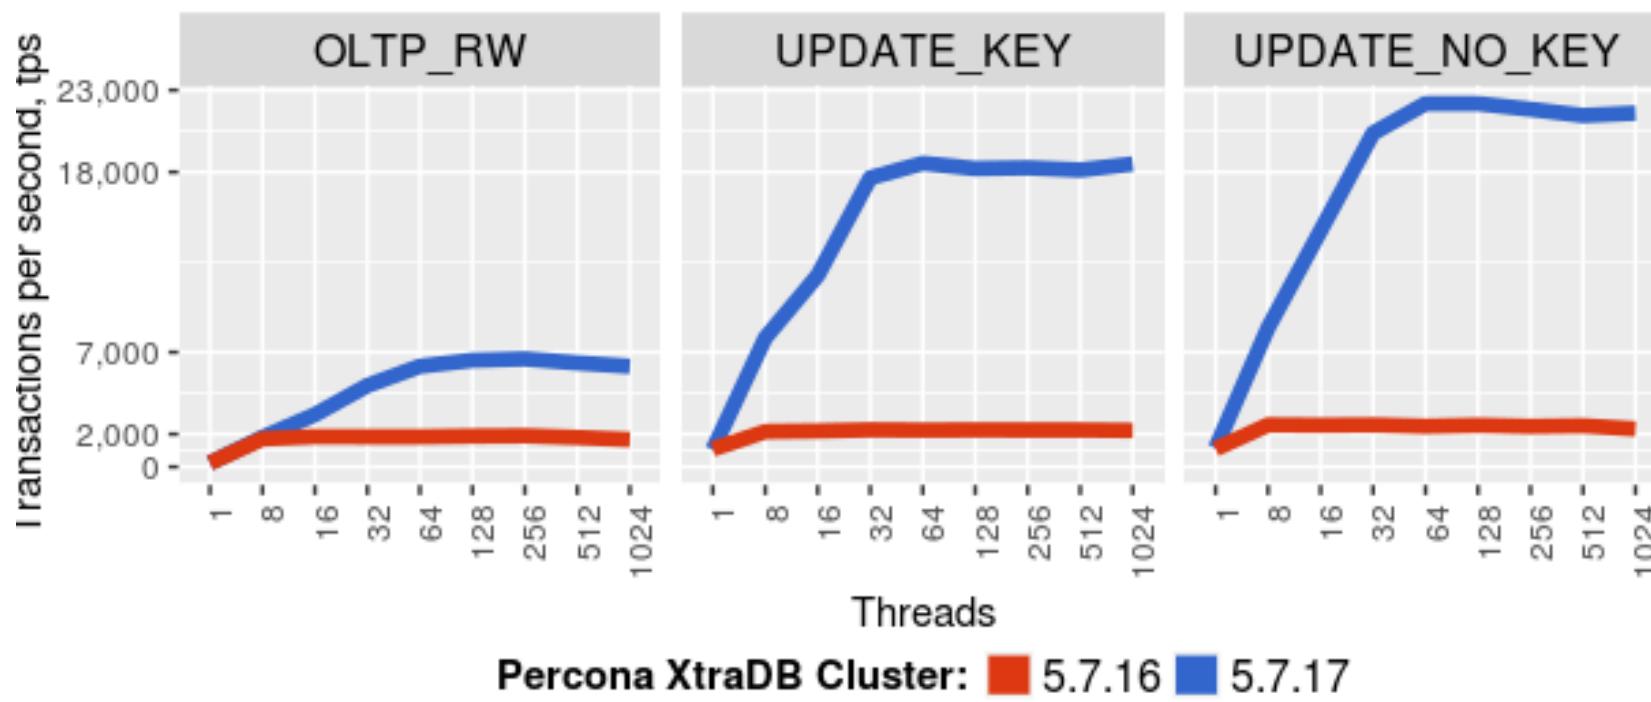
Transaction Commit Flow



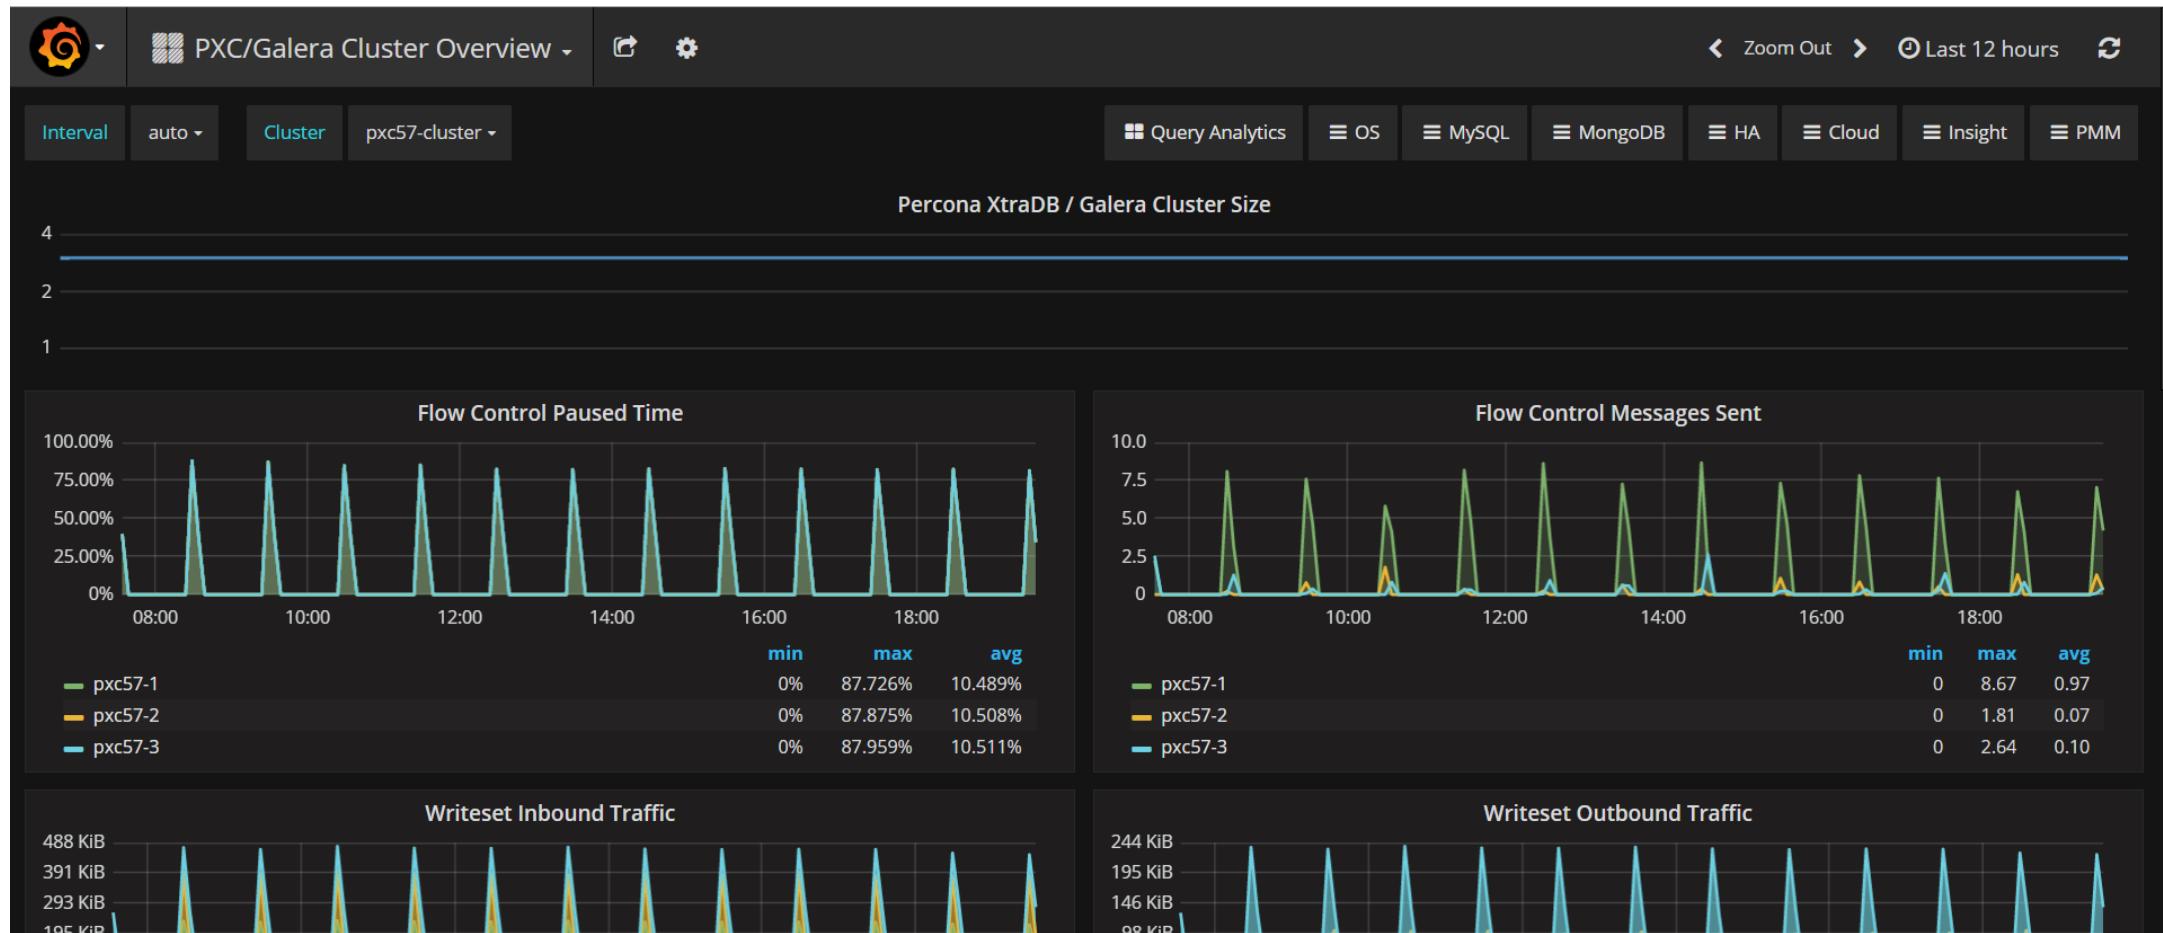
PXC 5.7 Performance Improvements

<http://bit.ly/2qGCr0T> and <http://bit.ly/2pzvAIW>

Sysbench: dataset 100tables/4M rows(100GB)
innodb_buffer_pool=150GB,innodb_doublewrite=1,
innodb_flush_log_at_trx_commit=1,sync_binlog=1
Box: 28 Cores+HT



PMM Dashboard for PXC



MySQL Cluster

Mostly In Memory Storage

Synchronous Replication

Pessimistic Locking

Conflict Detection with Async Replication

Niche Use Only

Sharding in MySQL

Typically Painful or Very Painful

No Built-In or Standard Solution

Many Custom App-Layer Implementation Exists

Thinking about Sharding

Query Routing to the right Shard

Data Aggregation for Cross-Shard Queries

Shard Balancing

Vitess

Stated by team in YouTube

Now Open Source in CNCF

Deployed at Slack

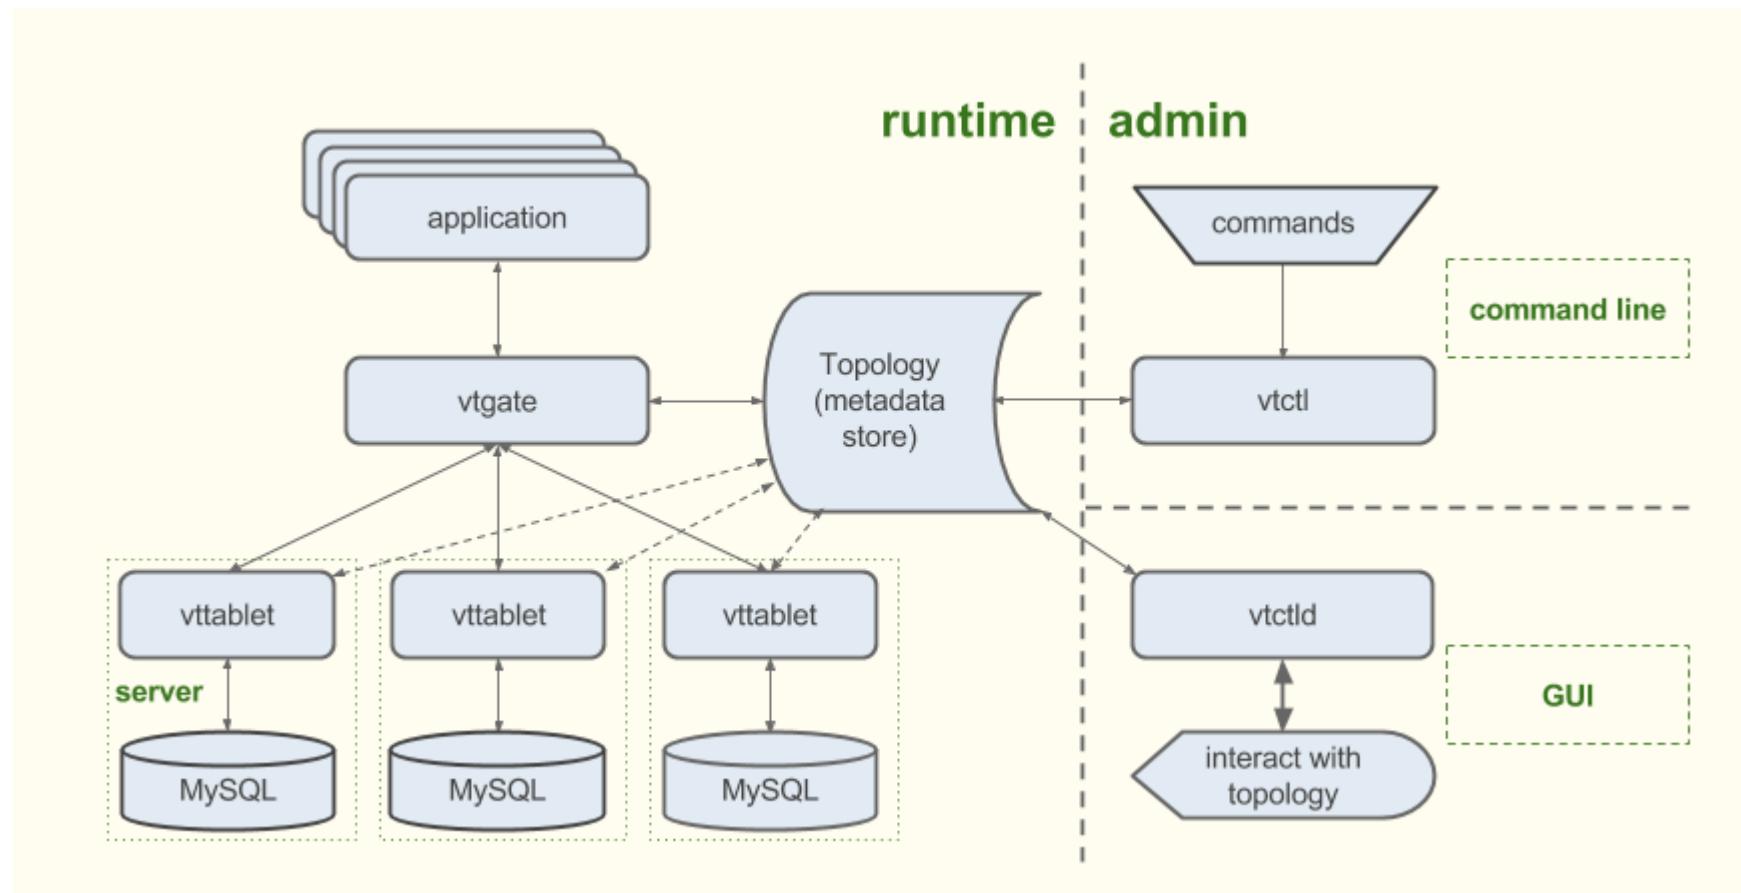
Supports MySQL Protocol

Supports Cross Shard Queries

Supports Re-Sharding



Vitess Architecture



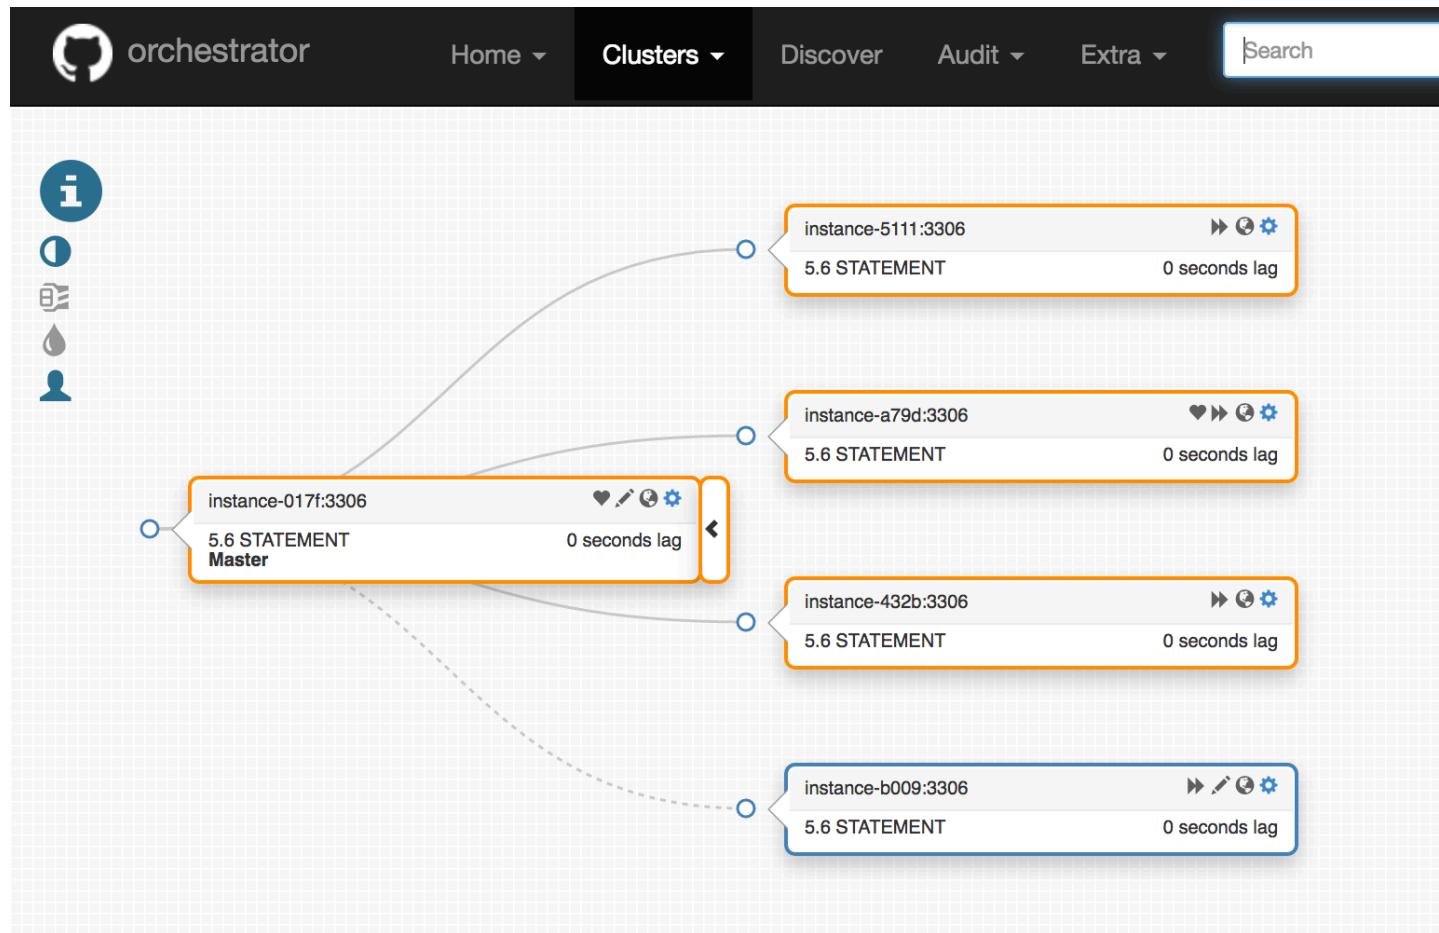
Failover in MySQL

MHA

MySQL Failover

Orchestrator

Orchestrator Screenshot



Traffic Management Solutions

DNS

- Not Immediate, Does not handle existing connections

HAProxy

- TCP/IP Port level. Does not understand MySQL Protocol

Elastic Load Balancer (etc)

- Like HA Proxy in the Cloud

MySQL Router

- Currently very basic

MaxScale, ScaleArc

- Proprietary Solutions

ProxySQL

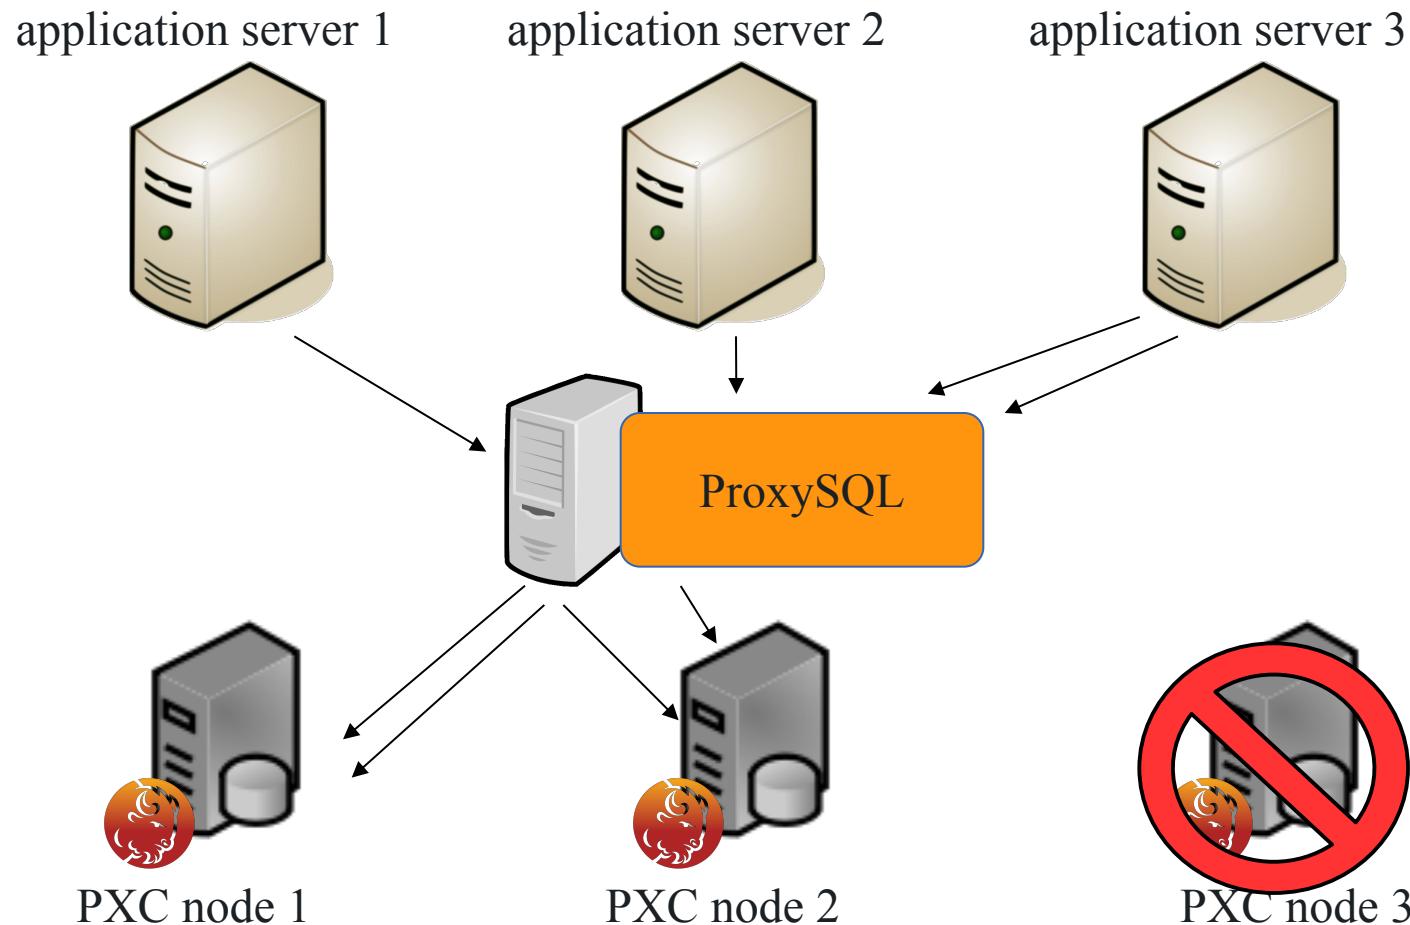
- MySQL Protocol Aware 100% Open Source Proxy Solution for MySQL

ProxySQL

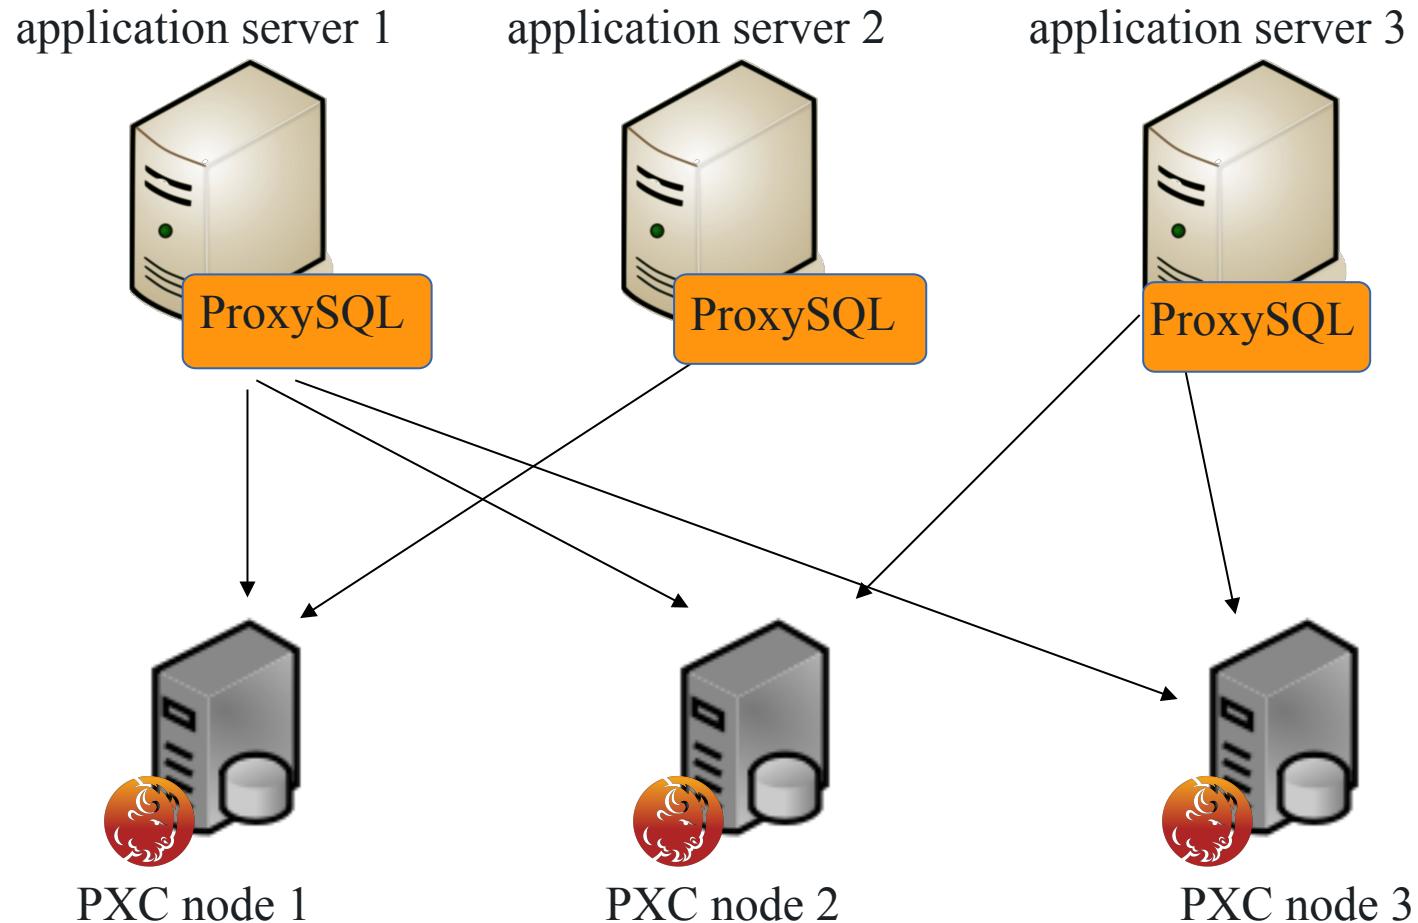
- 100% OpenSource
- Well Supported
- Read-Write Splitting
- Caching
- Query Rewrites
- Load Balancing
- 10k+ Connection Support



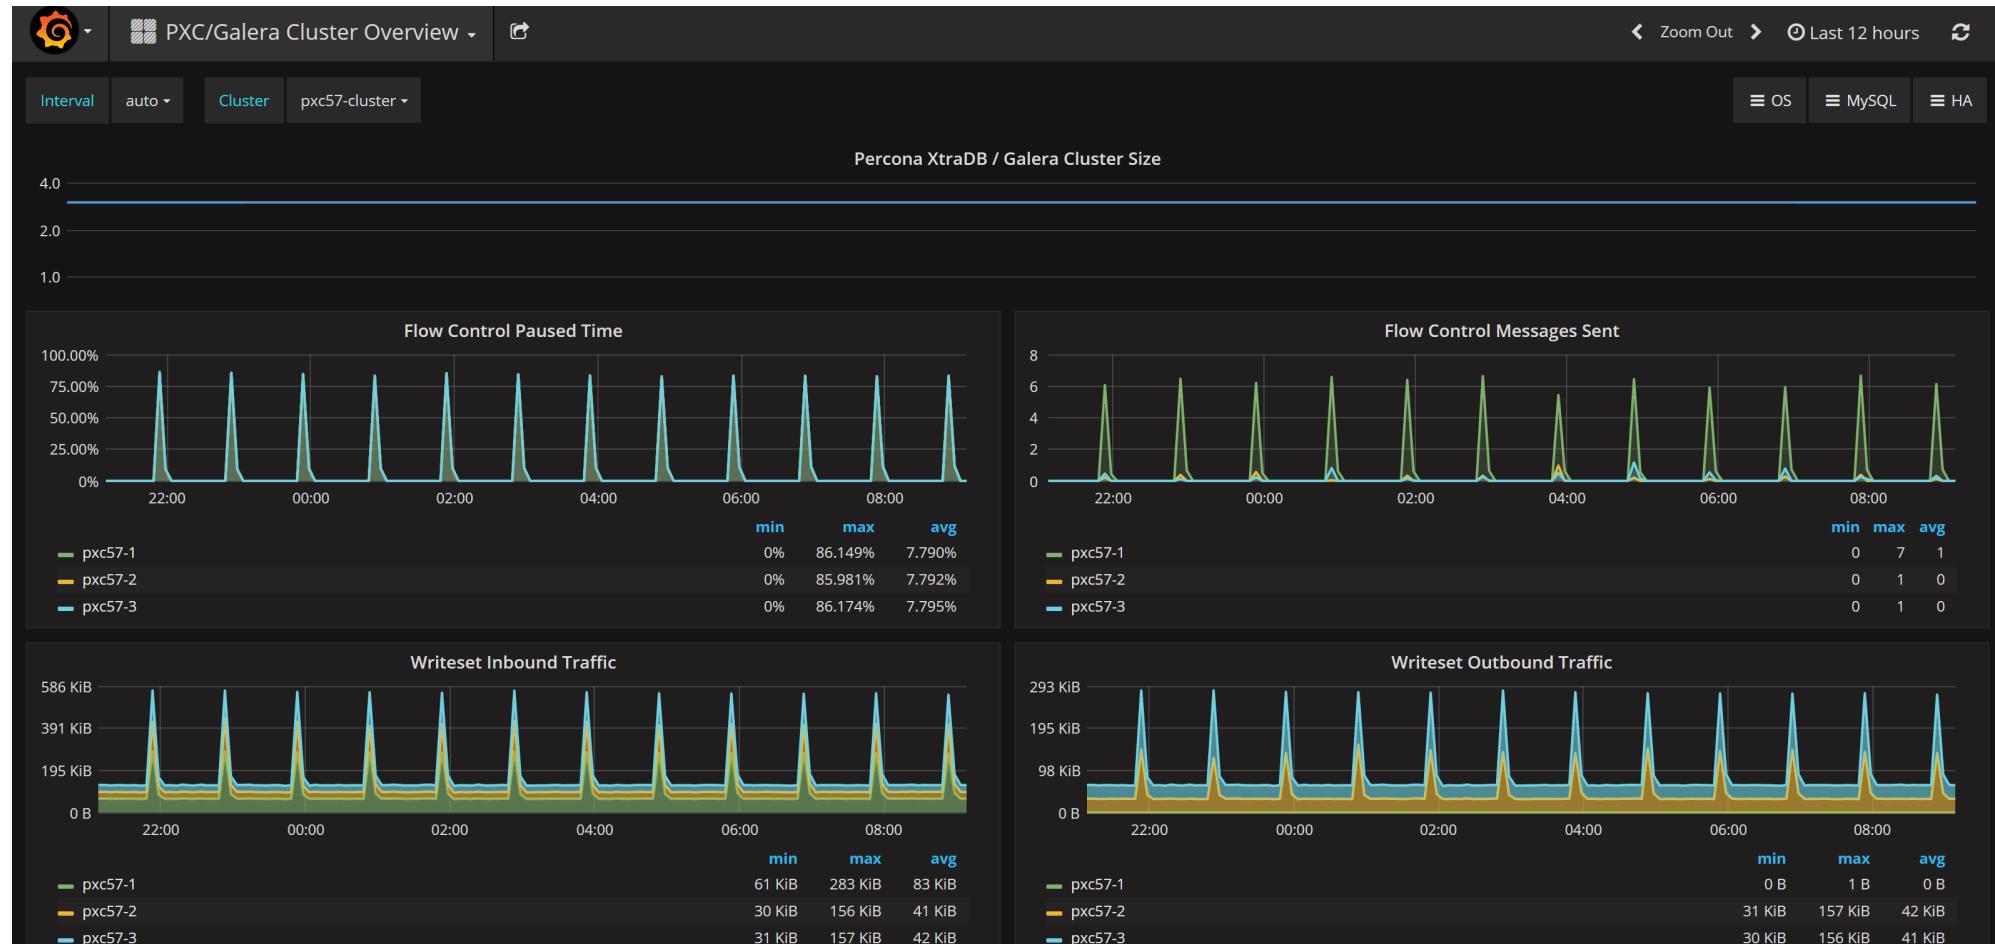
Dedicated shared ProxySQL



ProxySQL on application side



ProxySQL at PMM



Cloud Only Options

DBaaS in
the cloud
comes with
built-in HA
Options

- Amazon RDS
- Amazon Aurora
- Google Cloud SQL
- Azure Database for MySQL

Beyond MySQL

MySQL is a Part of your Data Infrastructure

MySQL is Good for Many
Things but MySQL is not
Good for Everything

MySQL can be used together with

Caching

- Redis,
Memcache

Full Text Search

- Elastic, Solr

Queue/Stream Processing

- Kafka

Analytics at Scale

- Hadoop,
Clickhouse

NewSQL ?

PingCAP TiDB

CockroachDB

YugaByte

Cross Technology Replication

Trigger Based Replication Tools

- SymmetricDS, Self-Made Solution

Kafka Gateway

- Debezium, Yelp's MySQLStreamer

Tungsten Replicator

- Replication from MySQL to Variety Data Stores

Custom Binary Log Consumers

- MySQL to Tarantool Replication
- MySQL to ClickHouse Replication

Summary

Lots of Options Exist for Building
Distributed Architectures with MySQL

Lots of Option for Integrating MySQL
with other Open Source Data Stores

SAVE THE DATE



PERCONA LIVE EUROPE FRANKFURT

5-7 November 2018
Radisson Blu
Frankfurt, Germany

Thank You!
